

## Authority, Categorical Reasons and Paternalism

**Abstract.** Instrumentalist accounts of authority justify authority on the basis of some good, benefit or service that authority relations provide for those subject to such relations. These accounts of authority are often criticised for justifying too much authority, and in particular, for justifying forms of authority which appear paternalistic: some authority relations do not seem justified even though they are, or would be, clearly beneficial to those subject to them. In this paper, I defend instrumentalist accounts of authority against this objection. I focus, in particular, on Joseph Raz's well-known and influential formulation of instrumentalism about authority, as articulated in his service conception, and on a version of the objection formulated by Stephen Darwall. I begin by arguing that, contrary to the suggestions made by several philosophers, instrumentalist justifications of authority would not be more plausible if restricted in scope to cases involving pre-existing obligations, or some particularly weighty subset of moral or “categorical” concerns. I suggest, instead, that worries about paternalism can be accommodated by instrumentalist accounts of authority in other ways.

### 1. Introduction

Authority, in the sense that I will be interested in here, is a distinctive kind of social relation, involving one person (or institution) who claims the ability to obligate another via command. Authority relations, understood in this way, are a ubiquitous aspect of our social lives: we find people who claim the ability to obligate others via command in politics and families, in classrooms and courtrooms, in militaries and sports teams. It is common to distinguish between *de facto* and *justified* authority. Given the distinctive normative claim at the heart of authority relations – the claim that the subject to a command *must*, or is *obligated*, to do as they have been told – authority relations are justified only where this claim can be justified, that is, when the subject to such a command really must, or is obligated to, do as they have been told. Someone has *de facto* authority when they are believed to have, or are recognised as having, justified authority.

Instrumentalist accounts of authority justify authority relations on the basis of some good, benefit or service that such relations help bring about. The most plausible instrumentalist justifications of authority restrict the kind of good or benefit that can contribute to the justification of authority relations by specifying that only goods or benefits which accrue to

those who are subject to the authority in question count. Establishing that relations of authority are beneficial for those who wield authority, or to some third party that subjects have no reason to benefit or assist, does not help to justify authority relations. If I claim authority over you, I must be able to explain why *you*, in particular, have a reason (indeed an obligation) to obey me. This does not mean, however, that only those authority relations which serve the subject's own interests, narrowly construed, are justifiable. A commander might direct a soldier to perform a dangerous manoeuvre; a politician might oblige a citizen to make tax contributions to a welfare scheme she will not herself benefit from. If authoritative directives like these can be justified on instrumentalist grounds, this is most readily explained by the fact that subjects to such directives have pre-existing reasons to contribute to the collective war effort, even if it is risky, or to assist other people, even at some cost to themselves. Authorities can be good for, or serve, their subjects in this broader sense, by helping them to comply with their pre-existing reasons, including reasons to help others, even if obedience to directives is not always in the subject's own best interest.

Instrumentalist justifications of authority are attractive in part because of what they imply about the value of authority. Relations of authority are relations of command and obedience: if I have (normative or legitimate) authority over you, I have the power to decide how you should act by issuing commands to you. The sorts of authority relations that dominate social life are paradigmatically asymmetrical and hierarchical: politicians and parents, bosses and commanders, judges and teachers, have the power to decide how their subjects should act, a power that their subjects lack. Asymmetrical power relations of this kind, it seems, are not intrinsically valuable: it is not a good thing, in and of itself, that some people have the normative power to tell other people what to do. We might even think of authority relations as intrinsically *bad*, since they appear presumptively problematic from the point of view of our freedom and equality. Other things equal, we might prefer a world with fewer and less far-reaching authority relations, than one with more. If authority relations are not intrinsically valuable, the most plausible justification for them will take the form of an instrumentalist justification: authority relations are justified only as a means to some worthwhile end.

The most sophisticated and influential instrumentalist account of authority in the philosophical literature is Joseph Raz's "service conception" of authority (1986, 1989). On this account, authorities are legitimate only if they provide their subjects with the service of improved reasons-compliance. In other words, the good that practical authorities provide their subjects – the same good which ultimately justifies relations of command and obedience – is the good of helping people better comply with the reasons that already apply to them. At the

centre of Raz’s service conception is the Normal Justification Thesis, which holds that subjects have an obligation (understood as a protected reason) to do as authorities direct them to do when they would better conform with the reasons that apply to them independently of the authority in question, if they treat the directive as binding and act on it, as opposed to acting on their own best judgment.

A common objection to accounts of authority like Raz’s is that they justify *too much* authority. This objection can be understood as consisting of two distinct but interrelated worries. First, people have worried that instrumentalist accounts of authority do not provide subjects with obligations to obey of the right kind. Many people think that our obligations to obey authoritative directives have the force of ‘moral,’ ‘categorical’ or ‘second-personal’ reasons, and worry that instrumentalist accounts of authority do not provide subjects with such reasons. Relatedly, given that instrumentalist accounts of authority do not seem to yield obligations to obey which meet this bar, many people have worried that they overgenerate, and pave the way for objectionable forms of paternalism. As we have just seen, any plausible instrumentalist account of authority limits the kind of good or benefit that counts toward the justification of authority relations by specifying that only goods or benefits that accrue to subjects of authority count. Though these goods or benefits are not *limited* to securing what is in the best interest of subjects, narrowly construed (authorities can be legitimate because they help their subjects be as altruistic as they have reason to be), instrumentalist accounts of authority do not tend to *exclude* justifications based on the self-interest of subjects. But authority relations which are justified only by appeal to the subject’s own best interest will strike many as paternalistic. If the only reason I have to obey you is that it would be better for me if I did, am I really bound to do what you tell me to?

In this paper, I defend instrumentalist accounts of authority against this objection. I introduce instrumentalism about authority by giving an account of Raz’s sophisticated version of it, in the form of the service conception of authority. I then present one of Darwall’s well-known objections against Raz, which takes the form of a putative counterexample involving a financial expert who decides to take command of my investment decisions. If the financial expert starts issuing directives to me, it seems as if her directives could satisfy all the conditions of Raz’s service conception. Nevertheless, intuitively, the financial expert lacks legitimate practical authority over me. Darwall’s objection, and others like it, seems to be motivated by the worry that the service conception does not yield obligations to obey of the right kind. Relatedly, the objection raises the worry that the service conception might not be sufficiently attuned to worries about paternalism. Unlike several other well-known criticisms against Raz’s

service conception, this objection has not yet received a satisfactory discussion in the literature.<sup>1</sup> Furthermore, reflection on it has important upshots for justifications of practical authority in general, as well as for wider philosophical debates about paternalism and perfectionism.

My argument goes in two stages. In the first part, I show why a common and initially attractive solution to Darwall's objection against Raz must be rejected. This solution holds that we should restrict the service conception of authority, and instrumentalist justifications of authority more generally, to a certain sub-set of particularly important or weighty moral concerns – such as cases where authorities help us improve our compliance with our pre-existing moral obligations, or categorical reasons, rather than reasons in general. I argue that any instrumentalist account of authority which is limited in this way would not yield a more plausible account of practical authority, since many paradigmatic cases of practical authority don't involve any such reasons. In fact, I argue that the widespread assumption about legitimate authority on which the proposed solution appears to be based – namely, that authoritative

---

<sup>1</sup> Most of the extant criticisms of the service conception can be put into three broad categories. First, many people worry that the service conception cannot take procedural considerations into account (the “proceduralist objection”) (e.g. Shapiro 2002). Second, many have worried that the service conception fails to capture the fact that authority is a relation between persons, rather than a relation between the subject and reasons (e.g. Coleman 2009). One strand of this objection holds that the service conception can't explain why authority is a relation between *specific* people or states and subjects: that children are subject to the authority of their own parents, rather than anyone who might help them improve their compliance with reasons, and that political subjects have an obligation to obey the directives of their own states, rather than the directives of some equally (or significantly more) competent state (the “specificity objection”) (e.g. Hershovitz 2011, p. 11-12). Third, some people have worried that the service conception justifies *too much* authority, by granting practical authority to people who, intuitively, have none. This paper is concerned with one version of this objection, which is related to the worry that the service conception might justify forms of authority which appear blatantly paternalistic. On a second version of this third objection, people worry that the service conception could grant legitimate authority to blatantly oppressive or unjust authorities. (See e.g. David Estlund 2009, p. 118, Christiano 2004, p. 278-279, and Perry 2013 for worries of this kind). Most of these objections to the service conception have already received lengthy discussions elsewhere. Daniel Viehoff has provided particularly careful and insightful discussions of many of these objections – though not of the objection that is my focus here. For a discussion of the first objection, see Viehoff (2011); for a discussion of the first version of the third objection, see Viehoff (2016).

directives must have the force of ‘categorical’ or ‘moral’ reasons in order to be legitimate – is mistaken.

In the second part, I return to Darwall’s objection. If limiting the scope of instrumentalist justifications of authority in the way suggested is not a viable option, how else might we respond to the worry that instrumentalist accounts of authority pave the way for objectionable forms of authority, including objectionably paternalistic directives? Here, my strategy is two-fold. First, I argue that instrumentalist accounts of authority can readily discharge most complaints about paternalism; and that any worries that remain do not undermine practical authority in the sense that I am interested in. Second, I explain why we have good reason to accept this, relatively undemanding, account of what practical authority amounts to.

The paper is organised as follows. In Section 2, I introduce instrumentalism about authority via Raz’s service conception of authority, and explain Darwall’s objection to such accounts of authority. In Section 3, I argue that a common response to Darwall’s objection, which is meant to rescue instrumentalist accounts of authority from justifying ‘too much’ authority, must be rejected. In Section 4, I discuss the remaining force of Darwall-style objections against instrumentalist accounts of authority. In Section 5, I conclude.

## 2. Instrumentalism about authority and the “too-much authority” objection

Any attempt at justifying authority on instrumentalist grounds faces two immediate challenges. Authoritative directives are meant to settle the question of how those subject to them are going to act, rather than simply provide subjects with additional reasons for action. In other words, authoritative directives, when successful, are in some sense binding or obligatory for subjects. This creates the first challenge for those who attempt to justify authority on instrumentalist grounds. Showing that some good or benefit (understood broadly) to the subject will come of obeying the directives of a putative authority is normally insufficient for showing that she has an obligation, or a binding reason, to obey. When and why does the good of following authoritative directives give subjects not just a reason, but an obligation, to obey?

A second challenge stems from the fact that we care not just about getting things right, but also about who gets to decide what to do. Sometimes, we’re willing to suffer worse outcomes in exchange for the freedom of being able to choose for ourselves. Since authority relations pose a threat to our freedom to make our own decisions, any plausible instrumentalist

account of authority must have something to say about how to balance the potential instrumental benefits of authority relations with the value of making decisions for ourselves.<sup>2</sup>

Raz's service-conception of authority provides well-known answers to both of these challenges. Raz, like most other theorists of authority, takes it as a starting point that legitimate authority entails obligation: if a legitimate authority issues a directive, subjects are bound, or have an obligation to, comply with that directive. An obligation, for Raz, is a *protected* reason for action: it is a reason which tells people what to do, while excluding, replacing or pre-empting other reasons for action. It is the pre-emptive nature of legitimate authoritative directives that explains the sense in which they are binding (1986, p. 60). On Raz's analysis, a *protected* reason is a combination of a first-order reason to  $\varphi$  and an exclusionary reason not to  $\neg\varphi$  for certain conflicting first-order reasons (1979, p. 18).<sup>3</sup>

This is followed by an account of how authorities are meant to function. Authorities, Raz thinks, should serve those subject to them, and more specifically, they should serve their subjects by helping them better comply with the reasons they already have, independently of the authoritative directive(s) in question. Well-functioning authorities, Raz proposes, are those which take the balance of these reasons – the *dependent* reasons, on his terminology – into consideration when issuing directives. Against the background of this account of what authoritative directives are, and how authorities ought to function, Raz proposes his Normal

---

<sup>2</sup> Many philosophers have thought that instrumentalist accounts of authority must also address a third challenge. Just like the value of *subjects* making decisions in a certain way (e.g. by themselves) can sometimes outweigh the value of getting things right, the value of *authorities* following certain decision-making procedures (such as making decisions democratically) might sometimes outweigh the value of getting things right. Raz's general account of practical authority doesn't say much about this third challenge. In part, I think this is because he aims to offer a unified account of practical authority, and the value of authorities making decisions in a particular way appears to differ across different contexts. For example, democratic procedures seem important in politics in a way that they do not in the context of parents making decisions on their children's behalf.

<sup>3</sup> In *The Morality of Freedom*, Raz tends to refer to authoritative directives as giving people “pre-emptive” reasons, remaining non-committal about how the phenomenon of pre-emption is best analysed. Raz's favoured explanation of pre-emption goes via “protected” reasons, and hence via exclusionary reasons. However, it is possible to explain the phenomenon of pre-emption, which appears central to the way in which authoritative directives function in practical deliberation, in other ways. Even those who are sceptical of the coherence of Raz's idea of exclusionary reasons, then, can accept that authoritative directives have a certain pre-emptive force.

Justification Thesis (NJT) – which specifies the ‘normal’ (but not the only) way to justify authoritative directives. This thesis, in response to the first challenge, explains how authoritative directives can give subjects obligations, and not just reasons, to comply:

**Normal Justification Thesis.** “The normal way to establish that a person has authority over another person involves showing that the alleged subject is likely better to comply with reasons which apply to him (other than the alleged authoritative directives) if he accepts the directives of the alleged authority as authoritatively binding and tries to follow them, rather than by trying to follow the reasons which apply to him directly.” (Raz 1986, p. 53)

Meeting the NJT, however, is neither necessary nor sufficient to establish practical authority. It is not necessary since there are other ways of establishing practical authority without meeting the NJT – for example, through consent. Meeting the NJT is also not sufficient for practical authority, since in order to establish that someone has practical authority, one additionally needs to show that there are no defeating contrary reasons for the establishment of such authority.<sup>4</sup> Raz mentions, in particular, that there is a need to satisfy what he calls the *independence condition*. This condition, in response to the second challenge, explains how we should balance the importance of getting things right with the importance of making our own decisions:

**The Independence Condition.** The matter (over which someone is said to have authority) is not one on which it is more important that people should decide for themselves than that they should decide correctly. (Raz 1989, p. 1180)<sup>5</sup>

---

<sup>4</sup> In *The Morality of Freedom*, Raz puts it in the following way: “... a complete justification of authority has to do more than to provide valid reasons for its acceptance. It has also to establish that there are no reasons against its acceptance which defeat the reasons for the authority” (1986, p. 56). In a later article, Raz formulates it slightly differently: “One needs also to show that there are no defeating contrary reasons.” (1989, p. 1180). The latter formulation is broader, since it could plausibly take into account not only reasons against *accepting* authority, but also reasons against *issuing* authoritative directives, or establishing authority relations. (I return to this point in Section 5.)

<sup>5</sup> In the cited article, Raz calls this the ‘autonomy condition’. However, in later work, Raz refers to it as the ‘independence condition’. The exact formulation of the condition varies. In the 1989 piece, Raz speaks of it being more important that ‘people’ should decide for themselves, whereas he later speaks of it as being more important that ‘one’ should decide for oneself (Raz 2009, p. 137). The two different

If both the normal justification thesis and the independence condition are fulfilled, “then, in general, the alleged authority is legitimate” (Raz 1989, p. 1181). The satisfaction of both the NJT and the independence condition thus at least approaches sufficiency.<sup>6</sup>

In many ways, the service conception of authority justifies very little authority. In particular, when applied to the political context, it denies that there is any general obligation to obey the law, as it establishes only “piecemeal” authority: whether or not any particular political directive, or set of directives, is binding is determined by asking whether any given individual has sufficient reason to obey that particular kind of directive. The service conception thus holds that subjects will have an obligation to obey some of the state’s directives, but not others; and sometimes, the very same directive can be binding for some, but not others. To use one of Raz’s own examples, an expert pharmacologist may not be subject to the authority of the government in matters of the safety of drugs since she likely won’t improve her reasons-compliance in this area if she follows the government’s commands rather than following her own judgments (1986, p. 74).

In spite of this, the service conception has been widely criticised for justifying *too much* authority, and in particular, for justifying authority in cases where, though the putative authority indeed “serves” the subject, intuitively, they lack legitimate authority. We find one well-known version of this objection in the writings of Stephen Darwall.<sup>7</sup> Darwall argues that

---

formulations – the first suggesting the independence condition applies on a group level, and the second that it applies on an individual level – may yield different answers about the scope of legitimate authority in some cases.

<sup>6</sup> Though Raz is uncomfortable with the idea that it always suffices for legitimacy, partly because he is aware that there might be other kinds of defeating reasons against establishing authority relations, not captured by the Independence Condition. For example, one class of further defeating reason that he mentions is the following: “One recurring kind of reason against accepting the authority of one person or institution is that there is another person or institution with a better claim to be recognized as an authority. The claim of the second is a reason against accepting the claim of the first only when the two authorities are incompatible, as are the claims of two governments to be legitimate governments of one country. Sometimes there are two compatible authorities whose powers overlap, as is the case with the authority of both parents over their children” (Raz 1986, p. 57).

<sup>7</sup> Jonathan Quong offers a similar example, involving a tour company which knows better than I do where I should go on holiday, where it is stipulated that I would comply better with the reasons that apply to me if I treat the tour company’s directives as binding, rather than act on my own best judgment (2010, p. 108-109).

merely securing better compliance with reasons isn't enough to ground the normative power of authorities to impose obligations. He gives the following example of a case in which an authority's directives would increase a subject's compliance with reasons, but where this intuitively seems insufficient to establish authority:

**Saving for Retirement.** I have prudential reasons to provide for my retirement that are independent of any obligations I might have, say, to provide for others whom I am answerable for supporting or of any obligation to support myself. It seems obvious that I do have such reasons and, moreover, that however important or valuable it might be for me to make my own choices, this latter value might not override (or sufficiently inform) the prudential reasons so that it could indeed make sense for me in prudential terms to put myself in the hands of a financial expert and simply follow her directives. Suppose, then, that I would better comply with the relevant prudential reasons if I were to do so and that there are no other reasons, or at any rate no sufficient reasons, for me not to do so. (Darwall 2013, p. 152-153)

Darwall argues that the service conception implies that the financial expert – we can call her Ruth, in accordance with Raz's own similar example from *The Morality of Freedom* (1986, p. 64) – has the power to impose obligations on me, since I would better comply with the reasons that apply to me if I accept Ruth's directives as authoritatively binding and simply follow them, rather than trusting my own judgment. This, he thinks, is clearly the wrong result: Ruth *doesn't* have legitimate authority over me. In Darwall's own terminology, improved compliance with non-'second-personal' reasons provides reasons of the wrong kind to ground the ability to impose moral obligations (2013, p. 148).

### 3. Limiting the scope of instrumentalist accounts of authority?

In response to objections like Darwall's, it might be suggested that instrumentalist accounts of authority, like Raz's service conception, must be restricted in scope if they are to provide plausible accounts of authority. Indeed, several philosophers have suggested (more weakly) that the service conception would be more plausible if restricted in this way, or (more strongly) that the service conception might be accepted as an account of legitimate authority if it is so restricted. Quong, for example, suggests that practical authority can be justified on instrumentalist grounds only where authorities improve our compliance with pre-existing

moral *duties* (2010, p. 116).<sup>8</sup> In a similar vein (though it is pitched as a defence, rather than a criticism, of the service conception), Ezequiel Monti (2018) has argued that, on the most plausible interpretation, the normal justification thesis must be limited in scope to cases involving compliance with “categorical” reasons – that is, reasons whose weight and stringency does not depend on the agent’s goals, tastes or desires.<sup>9</sup> Why? We might think of obligations not just as protected reasons for action, but as *categorical* protected reasons.<sup>10</sup> And, as Monti points out, if we think of obligations in this way, it seems as if the normal justification thesis could justify obligations if and only if the dependent reasons on which the authority bases their directives are themselves categorical (2018, p. 441).

This strategy for rescuing instrumentalist accounts of authority might appear to have some initial appeal.<sup>11</sup> Certainly, both Quong’s and Monti’s suggestions seem to yield the right

---

<sup>8</sup> See also Tadros (2020, p. 70), for the suggestion that the service conception would be more plausible if limited in scope in this way. However, Tadros (unlike Quong) ultimately thinks we should reject the service conception, even if it were limited in scope in this way.

<sup>9</sup> For a definition of categorical reasons along these lines, see e.g. Raz 2001, p. 168. Raz sometimes describes categorical reasons simply as reasons which do not “depend on the agents’ own goals” (see e.g. 1977, p. 224). In later work, Raz describes categorical reasons as reasons which apply to everyone with equal weight and stringency independently of people’s goals, tastes or desires (see e.g. 2001, p. 168).

<sup>10</sup> Raz himself often talks obligations in this way. For example, in *The Authority of Law*, Raz refers to his “general analysis of obligation” as holding that an obligation is a “categorical protected reason” (1979, p. 235). This isn’t the only place in which Raz speaks of obligations (and duties – Raz does not tend to distinguish between these) in terms of categorical protected reasons. He also does so in “Promises and Obligation” (1977, p. 223), in “On Respect, Authority and Neutrality” (2010, p. 291), as well as in his “Is There a Reason to Keep a Promise?” (2022, p. 187).

<sup>11</sup> Monti’s suggestion also has some support in Raz’s own writings. As Monti notes, when discussing the example of John, a putative authority on Chinese cooking, Raz suggests that how I ought to treat John’s directives “depends on my goals” (1986, p. 64). Raz might here be read as suggesting that John lacks authority over me because his commands don’t provide me with categorical reasons for action. Nevertheless, the textual support for this interpretation is not decisive, and I think it should be rejected. First, Raz is not consistent about defining obligation in terms of categorical reasons. Most notably, he never mentions categorical reasons in *The Morality of Freedom*, including in the three chapters dedicated to developing his account of authority. When Raz speaks of obligations to obey authorities in these chapters, he seems to have in mind simply pre-emptive reasons for action. Second, the passage where Raz discusses John, the putative authority on Chinese cooking, and Ruth, the financial expert, is

result when applied to Darwall-style counterexamples. In such examples, there is no pre-existing duty to do as the authority directs, nor do there seem to be any categorical reasons at stake. I don't have a categorical reason to do as Ruth says, since whether or not I should submit to her authority would appear to depend on my goals. (For example, if investing in the stock market is one of my hobbies, it would seem I have less of a reason to submit to Ruth's authority. And, conversely, if I *did* have a categorical reason or a moral duty to maximise my returns, we might be more inclined to think of Ruth as a justified authority.) However, as I'll show, this strategy must nevertheless be rejected. Adopting this kind of restriction on the scope of an instrumentalist justification of authority yields *too little* authority, and this will be especially obvious when we turn our focus away from cases of political authority towards other forms of practical authority.

Consider, first, the case of parental authority. Many paradigmatic cases of parental authority don't seem to involve helping children improve their compliance with categorical reasons. Consider parents who tell their children to practice the trumpet or attend football practice. Such cases seem to be paradigmatic cases of legitimate parental authority. And yet, the reasons on which such commands are based need not be categorical: they often depend on the child's own tastes, preferences, and goals. Or so I will argue. For while there are several reasons to suspect that there are always categorical reasons lurking somewhere in the background of legitimate parental commands, none of these reasons hold up under scrutiny.

First, consider that parents tell their children to do all kinds of things: to be nice to their friends, to share things fairly with their siblings, to eat their greens, and so on. Children generally have categorical reasons to do all these things. So, perhaps parental authority is in some sense holistic, deriving its legitimacy from the fact that parents help their children comply

---

ambiguous. Raz's claim that how I ought to treat John's directives "depends on my goals" allows for two readings. On the first, Raz claims that John lacks practical authority over me *whether or not I have the relevant goal*; on the second, that John lacks practical authority over me *unless I have the relevant goal*. It seems to me as if Raz does not intend to rule out the possibility that John or Ruth *could* be practical authorities over me. For example, he says that "whether or not there is a complete justification for me to regard their advice or instructions as guides to my conduct in the way I regard a binding authoritative directive depends on my other goals" (1986, p. 65). Third, given that Raz leaves no doubt that he intends for the service conception to provide a unified account of practical authority, the fact that Monti's reading would make the service conception a poor fit for many familiar cases of authority (as I go on to argue) counts strongly against it.

with categorical reasons, and naturally extending to cases where compliance with non-categorical reasons is at stake. However, this explanation of why parents have legitimate authority to tell their children to go practice the trumpet, or to play football, won't hold up. If there exists a relatively clearly delineated sphere of activity in which an authority's commands can be known to fail to meet the relevant justification of authority, there is no legitimate authority in that sphere (as Raz's case of the expert pharmacologist, mentioned above, is meant to show).<sup>12</sup> Of course, unlike in Raz's expert pharmacologist case, we are here supposing that a child's overall compliance with reasons would be better (not worse) if she were to follow her parents' directives to play the trumpet or attend football practice. But the important point is this: if we can show that our justification of authority fails in a clearly delineated category of cases – such as cases which involve improved compliance with non-categorical reasons only – there is no legitimate authority in such cases. And we *can* show this, since cases which involve improved compliance with non-categorical reasons only would fail to yield binding directives, and hence legitimate authority.

A second suggestion is the following. Perhaps parents are justified in telling their kids to play the trumpet or go to football practice independently of their goals, preferences, or desires. Part of parental authority involves helping children respect and engage with things of value, and one way of doing this involves getting them to participate in valuable activities. If trumpet playing is a valuable activity, one might think *any* child has categorical reason to play the trumpet, whether she wants to or not. Hence the presumed parental authority rests on improving compliance with categorical reasons after all.<sup>13</sup> However, I think this suggestion is mistaken. We don't all have reason to go practice the trumpet just because trumpet playing is

---

<sup>12</sup> As Raz formulates it, “a person or body has authority regarding any domain if that person or body meets the conditions regarding that domain and there is *no proper part* of the domain regarding which the person or body can be known to fail the conditions” (Raz 2009, p. 149, my emphasis). This implies that the expert pharmacologist is not subject to the authority of the government in matters of the safety of drugs even if it were true that her compliance with reasons would be better *overall* if she were to take *all* of her government's directives as authoritatively binding, as opposed to *none* of them. Here, my claim is that the domain of improved compliance with non-categorical reasons is a ‘proper part’ within the domain of parental authority.

<sup>13</sup> The suggestion is borrowed from Raz's own thought. Raz writes that we “have reasons to facilitate the realisation of anything of value whose realisation is not deeply impossible, regardless of whether or not its realisation is one of our ends” (Raz 2011, p. 142). Such general reasons to respect things of value seem to apply to us all equally, independently of our goals. Hence, they are categorical.

a valuable activity. We should distinguish, as Raz does, between reasons to *respect* and *facilitate the realisation of* things of value, on the one hand, and reasons to *engage* in any particular valuable activity, on the other.<sup>14</sup> It seems plausible to suggest that we all have categorical reasons to recognise and acknowledge things of value, as well as categorical reasons to protect and preserve things of value, independently of our own goals and preferences. For example, even someone who doesn't particularly care for early Baroque painting has reason not to destroy a Caravaggio, and take good care of such a painting if it were to come into their possession. Such reasons plausibly apply to everyone with equal strength and stringency. However, someone who doesn't care for early Baroque painting has no reason to take a class on Baroque painting or to seek out the nearest museum featuring Baroque masterpieces. What we're dealing with in the case of a parent telling their child to practice the trumpet or attend football practice are reasons to *engage* in a particular activity of value, as opposed to reasons to respect or facilitate trumpet playing as a valuable activity. Whilst the latter plausibly offers categorical reasons, the former reasons are not categorical, since their weight and stringency depend on the agent's goals and preferences.

A third and final suggestion goes as follows. Whilst we might not have a categorical reason to play the trumpet just because trumpet playing is a valuable activity, perhaps we have a categorical reason to learn the skill of committing to a project and seeing it through. Put slightly differently, perhaps we all have categorical reason to develop our capacity for living an autonomous life, which involves the capacity to formulate worthwhile projects and see them through. (And even if we're unsure if all adults have such categorical reasons, it seems plausible that all children do, and that's all we need for the objection to be successful.) If so, that could explain the basis of legitimate parental authority in some of the cases I have just mentioned.

This line of reasoning may explain part of the reason why we are so confident in saying that there is legitimate parental authority involved in those cases where a parent tells their child to go practice the trumpet, or to attend football practice. An important part of any parent's job rests on ensuring that their children develop a capacity for autonomy, and that they learn how to formulate and stick to worthwhile goals. Nevertheless, I do not think the suggestion works. Even if there were no categorical reason of this kind at play in the example, the command could still be legitimate. To see this, consider a variation of our case in which the child needs no encouragement to practice the trumpet: they are fully committed to practicing on a regular

---

<sup>14</sup> Raz discusses this distinction in his *Value, Respect and Attachment* (2001, Ch. 4).

basis. However, they keep on repeating certain basic mistakes when practicing. A parent might, in such a case, step in to give the child additional instruction about how to play. Assuming that the parent in question actually knows what they are talking about, such instructions could be legitimate. In this case, it is clear that what is at stake isn't the child's capacity to formulate and stick to a worthwhile goal, but only their proficiency at playing their chosen instrument.

My argument, so far, has sought to establish that it cannot be a condition of legitimate practical authority, in general, that such commands are based on, or result in, obligations which are categorical in nature. I've appealed to cases of parental authority to show this idea to be mistaken. However, there are several reasons why we might think parental authority is special; reasons which might make us think that the argument just given doesn't extend to other forms of authority. Most obviously, parental authority is exercised over children, whose cognitive capacities are still developing. Furthermore, most people think parents have special obligations towards their children, and the existence of these obligations might muddy our intuitions about the source and the scope of parental authority. Finally, worries about justifying 'too much' authority, and especially paternalistic forms of authority, appear somewhat less pressing in the context of parental authority. Parental authority *should* (at least sometimes, and in some circumstances) be paternalistic.<sup>15</sup> So while it may seem mistaken to restrict the scope of an instrumentalist justification of authority to categorical reasons only in the case of parental authority, one might be sceptical that the argument generalises.

Reflection on a range of other cases of practical authority, however, reinforces the argument just made, and shows that the argument applies also to cases involving rationally competent adults. Consider, for example, the practical authority exercised by university teachers over their students. When I assign readings to my students, I am not merely advising them on what might be a good idea to read (as I might when a student asks me for additional reading recommendations). Nor am I threatening them with adverse consequences if they do not comply (I do not have the power to punish my students for not doing the reading). Instead, I am instructing them to do a particular reading, and I intend for my instructions to be binding, in the following sense: I intend for them to settle the question of what my students will read before the next class. My instructions are meant to pre-empt or exclude various other reasons

---

<sup>15</sup> This should not be taken to imply, however, that charges of paternalism are inappropriate, or cannot be raised, in the context of parental authority. Children (and especially teenagers) often accuse their own parents of paternalism; in some such cases, children are correctly pointing out a flaw in the way that their own parents treat them.

my students might have (e.g. to read other things instead). This, then, is a clear example of practical authority, understood in terms of one person's claim to having the ability to issue binding directives to another. Intuitively, on the assumption that my instructions are genuinely helpful for my students, my authority is also legitimate. Yet, my students' obligation to do as they've been told is not categorical. Its strength and stringency depends on how much they care about doing well in philosophy.

Granted, there is a possible confounding factor in this case of legitimate practical authority, not present in the case of legitimate parental authority: consent. It might seem plausible that the practical authority of university teachers rests on consent, rather than on any instrumentalist justification of authority along the lines of the service conception. If that's right, and we think that consent gives rise to categorical reasons, then there may be categorical reasons for my students to obey my directives.

The problem with this explanation, however, is that my students have not consented to obeying me. University education, of course, is voluntary (in the sense that it is optional: students can choose whether or not to go to university). But we cannot infer from the fact that some activity is voluntary in this sense (and morally permissible) that people have therefore given their valid and binding consent to it. My students are asked to consent to many things when they start university – e.g. to pay their course fees, and abide by the university's code of conduct – but they are not asked to consent to obeying their professors. Nor does their enrolment in my class constitute any kind of implicit consent to obeying me. To test this thought, consider a case in which I don't know what I am talking about, and instruct my students to read irrelevant texts that won't help them become better philosophers. In such a case, my students are under no obligation to do what I tell them to, based on the mere fact that they have enrolled in my class. (And they would certainly not be in breach of any of their contractual obligations if they were to disregard my directives or drop my class.) Implicit or explicit consent or contract, then, is not a plausible grounds of my legitimate authority over my students. The more plausible explanation of my authority over my students is an instrumentalist one, based on my expertise. My students have an obligation to do as I tell them to only if, and because, this will help them become better philosophers.

The argument just given shows, a fortiori, why Quong's suggestion – that all forms of legitimate practical authority should be limited to cases involving pre-existing *duties* – should be rejected. There are no such pre-existing duties in the paradigmatic cases of legitimate authority considered above. Relatedly, it should be clear that parents' or teachers' commands need not be based on any distinctively *moral* reasons which apply to children or students, on

any plausible understanding of what a moral reason amounts to. This shows that we should reject a widespread assumption about legitimate authoritative directives – namely, that such directives necessarily give their subjects a distinctively moral or categorical duty to comply. Often, legitimate parental commands are based on nothing but their child’s own interests or well-being. Similarly, the legitimate directives of teachers are often based on nothing but their students’ interests. Because of this, authoritative commands – even those of a paradigmatically legitimate kind – do not necessarily provide people with categorical or moral reasons for action. Intuitively, in many cases, people’s obligations to obey legitimate authoritative directives don’t have any distinctively moral or categorical force.

At this point, let me address two potential objections to the argument I have just given. First, the reader might be wondering in precisely what sense a child or a university student has an *obligation* to do as they have been told, in the cases I have been describing. It is a characteristic feature of obligations that they constrain our freedom, in some sense. Obligations, it is sometimes said, are meant to be ‘non-optional.’ But if obligations are non-optional, can there ever be a genuine obligation to obey authority, if the subject can escape her obligation to obey simply by changing her goals or preferences? Similarly, in what sense are my students really *obliged* to do what I tell them to, given that their participation in my class is voluntary, and they could easily escape my authority by dropping my class?

The answer is that the relationship between obligation and voluntariness is complicated. Obligations can constrain our freedom, in one sense, even when they are voluntary assumed and can be voluntarily escaped. In other words, there are different kinds of ‘bindingness,’ and not all kinds of obligations are binding in all senses. Consider obligations of friendship. These obligations are, paradigmatically, voluntary. We can choose who we want to be friends with, and we can end friendships we no longer want to be in. The obligations of friendship are also non-categorical: they do not apply to everyone with equal strength and stringency. (Obligations of friendship, then, seem to speak decisively against any account of obligation understood in terms of *categorical* protected reasons.) Yet, the obligations that arise out of friendship constrain us in many ways, presenting themselves to us as non-optional in our practical deliberation. Given that my best friend recently had a breakup and needs my support, I must answer when she calls, even though I’m in the middle of doing something else. If she breaks a leg and ends up in the hospital, I must go to visit her, even though the hospital is far away and I had already made other plans. And so on. The fact that we recognise obligations of friendship as *obligations*, and not just weighty reasons, shows that we recognise some obligations as constraining our freedom, even though they are voluntary.

But how do they constrain our freedom, exactly? The bindingness or non-optionality of obligations of friendship, as we have just seen, cannot be explained in terms of the fact that they cannot be escaped, nor in terms of their categorical force. Instead, their bindingness is best explained in terms of their pre-emptiveness. Obligations of friendship are meant to exclude or pre-empt other reasons we have for acting. They stop us from letting certain considerations influence our practical deliberation about what to do (or, alternatively, stop us from acting on certain kinds of reasons).<sup>16</sup> For example, given that I am obligated to go visit my friend in the hospital, I'm no longer free to simply do what I most feel like doing. This obligation does not necessarily preclude all other considerations that bear on my actions, however. Importantly, it need not preclude me from acting on the considerations that speak in favour of ending our friendship. (Though there may be independent constraints on how one may permissibly end a friendship. The fact that one no longer feels like going to visit one's friend in the hospital is not a good reason to end a friendship; hence, it seems impermissible to end a friendship *for this reason*, while it remains permissible to end it for other reasons.)

My suggestion is that in the cases considered above, there is an obligation – understood simply in terms of a protected reason – to do as one has been told. Protected reasons bind us in the sense that they exclude or pre-empt other reasons for action; thereby constraining what considerations we are allowed to act on, or to let into our practical deliberation. Such reasons need not be ‘all-things-considered’ reasons: protected reasons can, and sometimes do, conflict with the overall balance of reasons. This can happen, for example, where the overall balance of reasons speaks in favour of breaking a promise, or disobeying a command issued by a legitimate authority. Protected reasons also need not be ‘decisive’ reasons: obligations can conflict, and in such cases, a weightier obligation may speak decisively against conforming with another, less weighty obligation.

Understanding obligations in this way has the advantage of helping to explain the relationship between practical authority and voluntariness. Many forms of practical authority, including political and parental authority, are non-voluntary, in the sense that we cannot avoid these kinds of authority by saying no, opting out, or indeed by changing our desires, preferences or goals. But many other familiar forms of legitimate practical authority – such as the authority

---

<sup>16</sup> On Raz's account of pre-emption, pre-emption is explained by an appeal to exclusionary reasons. These are reasons *not to act* on certain reasons. Other accounts of pre-emption, such as that of H.L.A. Hart (1982, Chapter 10) instead stress the idea that pre-emptive reasons are reasons *not to deliberate* on certain reasons.

of university teachers, workplace bosses, and sports coaches – are voluntary. We can opt out of, or exit, such authority relations if we wish to do so.<sup>17</sup> This point holds even if we think that these forms of authority are based on consent, since consent to these kinds of authority relations can be revoked. (And presumably, it can be revoked at any moment, unless we have agreed to be bound by someone else's commands for some specific duration of time.)<sup>18</sup> Yet, this doesn't threaten the status of these kinds of authority relations as instances of practical authority. We must allow, then, that there can be authoritative directives which are binding, and yet optional for the subject. Understanding the bindingness of authoritative directives in terms of protected reasons explains how this is possible.

Let me turn to the second objection. Many people think that complying with authoritative directives is the sort of thing that we can be held morally accountable for, either by the authority figure themselves, or by third parties. But complying with non-categorical or non-moral reasons, it might seem, is not the kind of thing we can be held morally accountable for.<sup>19</sup> No one else will be harmed or disappointed if I fail to write the article I've been planning to write, or if I show up to tennis practice hungover. Similarly, if I fail to respect a place of great natural beauty, it might be that no one at all is disappointed or harmed by my action. In both cases, we might doubt whether I owed it to other people (or to anyone at all) to comply

---

<sup>17</sup> As already noted, the fact that these forms of practical authority are voluntary or optional for us shouldn't be confused with the idea that their legitimacy is grounded in our consent..

<sup>18</sup> This point also shows why it is wrong to associate categorical reasons with reasons that we 'can't (easily) escape,' and non-categorical reasons with reasons that we 'can (easily) escape.' An obligation to obey authority which is based on our consent to obey would presumably be categorical in nature, in that its strength and stringency would not depend on our preferences, goals or desires. Instead, it would depend on the fact that we had consented. Yet, if such an obligation can be escaped by revoking our consent, it is easy to escape such an obligation. Contrast this to a non-categorical obligation that applies to us in virtue of our own desires, preferences or goals. We may not be able to escape such obligations at will, simply because we do not necessarily have the power to change our desires and preferences at will. We can typically *shape* them through our voluntary actions, but this is a process which tends to take time, and typically requires much more than (the communication of) an act of will.

<sup>19</sup> This worry seems to play a big role in explaining e.g. Darwall's dissatisfaction with the service conception: Darwall thinks that we can't be held accountable for complying with what he thinks of as non-'second-personal' reasons. In general, this worry would loom large for anyone who takes a 'relational' approach to morality, on which all moral obligations are obligations owed *to* some other person.

with these reasons; and consequently, whether others can hold me ‘morally accountable’ for my failures. Nevertheless, it would be a mistake to conclude that authoritative directives could not be justified when they are based on these sorts of reasons. We have good reason to keep distinct questions of legitimate practical authority, on the one hand, and questions of whether the authority figure, or other people, have the moral standing to demand that I comply with any particular authoritative directive, on the other. Most obviously, many of the most familiar forms of practical authority are set up for the benefit of the commanded, rather than the commander.<sup>20</sup> In light of such cases, it seems that we must either conclude that many authoritative directives involve a form of bindingness which falls short of the bindingness of ‘moral’ or ‘other-regarding’ obligations, or else reject many familiar forms of practical authority as practical authority altogether.<sup>21</sup>

Before moving on, let me say something about the scope of my argument. While the case of parental authority and the authority of university teachers over their students provide two of the clearest counterexamples to the idea that authoritative directives are legitimate only if they help people better comply with their categorical reasons, my argument may be extended to further cases of practical authority as well – though admittedly, some of these further cases are more controversial, and my conclusion does not depend on them. Consider, for example, the authority of bosses in the workplace, or the authority of coaches or captains in sports teams. The strength of a worker’s obligation to obey her boss typically depends on how badly she needs or wants to keep her job. (Which is not to say that workers never have a moral reason or duty to obey their bosses. Sometimes they do, such as when the surgical assistant is obliged to

---

<sup>20</sup> It may be noted that this goes against one of Hobbes’s famous observations about the difference between counsel and command. Hobbes claims that counsel is directed at the benefit of counselled, whereas commands are directed towards the benefit of the commander (2008, p. 169). Hobbes thinks this is what explains why we are bound by command, but not by counsel: “a man may be obliged to do what he is commanded... but he cannot be obliged to do as he is counselled, because the hurt of not following it, is his own” (2008, p. 169). Needless to say, I am using the term ‘command’ in a sense that differs, in this respect, from Hobbes’s.

<sup>21</sup> Some readers might remain sceptical about defining obligations in terms of protected reasons, as I have done here. I hope to have shown why there is good reason to define obligations in this way. Nevertheless, it is not my main aim in this article to argue for this particular account of obligation, and my argument does not rest on the terminological question of whether this is the best definition of the term. The sceptical reader should feel free to replace my use of the term “obligation” with “protected reason.”

obey the surgeon.) The strength of a football player’s obligation to do the drills their coach gives them depends on their goal of becoming a better football player. Of course, consent is a possible confounding factor in these cases, too. And while I agree that consent, or a promise to obey, *could* ground practical authority in some of these cases, I am sceptical of the idea that consent is always, or even typically, the grounds of these forms of legitimate practical authority. It is a mistake to suggest that, whenever we join a sports team or a workplace, we implicitly consent to, or make a promise to, obey the coach’s or boss’s directives. When we join a sports team, we may implicitly consent to many things – to show up to practice, to pay the requisite fees, and (perhaps) to do our best – but none of these amount to consenting to obey the coach. In the employment case, what we have agreed to is usually spelled out in our employment contracts. Typically, we agree to perform a certain job or task, but we do not necessarily consent to or promise to obey the boss. (Though if obeying the boss is the best way of doing the job, we may indeed be under a duty to obey the boss, derived from our contractual duty to do the job.)<sup>22</sup>

Finally, consider cases of political authority. Though it might seem like restricting an instrumentalist account of authority to helping us improve our compliance with categorical reasons or moral duties is most plausible when we focus on cases of political authority, it is worth noting that, in this context too, such a restriction seems to justify too little authority. Many people are happy to concede that e.g. mandatory pensions schemes are (or could be, if designed in the right way) legitimate, even though the reasons that individuals have to make pensions contributions presumably depend in part on their preference for spending now versus saving for later. For those who agree, the lesson learned from reflecting on examples like Darwall’s cannot be that no political authority, or practical authority more generally, is legitimate unless based on reasons which apply to us with a certain categorical force.

#### 4. Authority and Paternalism

My argument, so far, suggests that what might initially appear to be a promising way of rescuing instrumentalist accounts of authority from the charge that they justify too much authority has unacceptable results. This, in a sense, brings us back to square one. If we don’t limit the scope of instrumentalist accounts of authority in this way, won’t our account of authority overgenerate cases of justified authority? Isn’t there something obviously

---

<sup>22</sup> See also Viehoff (ms.) for a similar argument.

paternalistic about granting the financial adviser practical authority over me? To address these questions, it will be necessary to first say a bit more about how the charge of paternalism should be understood. I argue, first, that instrumentalism about authority has no special problem discharging worries about paternalism. Second, while I acknowledge that we may have a paternalism-related complaint against authority relations like the one present in Darwall's counterexample, it is not clear why such complaints undermine the justification of authority.

There are three broad kinds of accounts of what makes an act paternalistic, which can be distinguished by how they characterise the wrong involved in paternalism. On the first type of account (“outcome-based” accounts), paternalistic policies are presumptively objectionable because they coerce people, restrict their liberty, or threaten their autonomy, in a way that is problematic or unjustified.<sup>23</sup> On the second type of account (“motive-based” accounts), paternalistic acts are presumptively objectionable because they involve some agent A aiming or intending to substitute her own judgment or agency for B's judgment or agency, in some sphere of action which legitimately lies within B's control.<sup>24</sup> On the third type of account (“expressive-content” accounts), paternalistic acts are presumptively objectionable because they express the idea that one person's judgment, or capacity for action, is better than another's, in a way that is disrespectful.<sup>25</sup>

If one accepts the first kind of account of paternalism, one needs to ask whether instrumentalist accounts of authority constitute an unjustified restriction or threat to our liberty or autonomy. Do they? Authority, in our sense, certainly involves a kind of restriction on subjects' liberty, in the sense that those who are subject to legitimate authority are, normatively speaking, bound to do as they have been told, and would do something wrong if they disobeyed. Any plausible instrumentalist account of authority, however, will contain a provision about how to balance the importance of making our own decisions, with the importance of getting things right. Raz's independence condition does just that: it states that authority is justified only in those cases where it is not more important that people should decide for themselves, than that they should decide correctly. Those who are worried that Raz's service conception paves the way for paternalistic forms of authority might think that Raz seems all too ready to trade the value of autonomy against the other potential benefits of authority relations. Alternatively, they might be worried that Raz's ideal of autonomy fails to capture some

---

<sup>23</sup> See e.g. Sunstein and Thaler (2003).

<sup>24</sup> See e.g. Shiffrin (2000).

<sup>25</sup> See e.g. Quong (2010, Chapter 3) and Cornell (2015).

important aspect of autonomy or freedom altogether.<sup>26</sup> But these are not objections against the service conception, or instrumentalist accounts of authority more generally, but objections against the substantive account of autonomy that informs the independence condition. If we think the value of individual autonomy in certain spheres of action – for example, when it comes to retirement savings – is so important that it could not be outweighed by the potential instrumental benefits of authority, we would simply reach a different verdict about whether or not there is, or could be, legitimate authority in putative counterexamples like Darwall’s. The same may be true if we use a different ideal of autonomy or freedom to inform the independence condition. The important point, however, is that instrumentalist conceptions of authority can in principle be combined with different substantive accounts of autonomy, yielding different interpretations of the independence condition.

Take motive-based accounts of paternalism, next. We can grant, from the start, that in the sorts of examples we are interested in, the putative authority figures might indeed have the relevant aim or motive: the putative authorities in question, we might suppose, are indeed aiming or intending to substitute her own judgment or agency for the subject’s judgment or agency. But does it do so within some sphere of action which legitimately lies within B’s control? The difficulty (as proponents of motive-based accounts of paternalism themselves note) is that legitimate authority intuitively affects what lies within someone’s legitimate sphere of control. For example, if my legitimate parental authority extends to being able to decide which dentist my child will go to, my decision to take her to some particular dentist is not objectionably paternalistic. Hence, a motive-based account of paternalism like Shiffrin’s cannot be used to mount a successful critique of authority as objectionably paternalistic prior

---

<sup>26</sup> These, at least, are two possible interpretations of the dispute between Raz and his critics. There might be others. For example, the critic might not be worried that Raz has misjudged the importance of autonomy, nor that he has failed to capture some important dimension of autonomy or freedom. The critic may instead think that, in some sphere of self-regarding choices, no one else has the capacity to change my normative situation via command – not because me making the choice autonomously (in some sense of that term) is more important than my getting it right, but rather because my self-regarding reasons for action are, simply ‘nobody else’s business.’ I address this alternative interpretation of the dispute between Raz and his critics below, in considering the idea that there may be some sphere of action in which I have a right to remain normatively unaffected by others’ authoritative commands. I thank an anonymous reviewer for this point.

to giving some independent specification of what constitutes someone's legitimate sphere of control.

But perhaps there *is* some such specification we can give of what constitutes someone's legitimate sphere of control, a specification that would both place limits on what could count as legitimate authority, and hence also on what actions count as paternalistic. There seems to be two possible ways of doing this. On the first approach, we might say that people's legitimate sphere of control consists of the sphere in which the value of them being able to make their own decisions outweighs other values, such as the potential instrumental benefits of being subject to someone else's authority. But this, again, sounds a lot like a restatement of the independence condition. If this is how we should spell out the contours of people's legitimate sphere of control, then, any instrumentalist conception of authority can avoid the charge of paternalism by including a provision like the independence condition.

The alternative, it seems, would involve abandoning the idea that the value of autonomy can be weighed against other values, and instead positing that there is some sphere of control in which we have a *right* to autonomy, or a right to make our own decisions, which cannot simply be outweighed by countervailing values or benefits. But how should such a right be understood? Here, we must tread carefully. On one of the most natural interpretations of what a right to autonomy, or a right to making one's own decisions, amounts to, this involves a right against certain forms of coercive interference in one's decision-making. But recall that legitimate practical authority, as I am understanding it here, doesn't necessarily imply anything about the legitimacy of coercion or enforcement. That means that legitimate practical authority need not threaten people's right to non-interference, in *this* sense. This might sound paradoxical, but it is simply a consequence of divorcing the analysis and justification of practical authority, understood in terms of the normative capacity to impose obligations, from the analysis and justification of coercion and other forms of interference with a subject's decision-making. There is nothing strange about thinking both that some subject is obliged to do as they have been told, at the same time as maintaining that it would be wrong (for the authority, or for anyone else) to interfere with their decision to obey or disobey.<sup>27</sup> Consider, again, the practical authority I have over my students. Such practical authority, it seems to me, is compatible with the idea that it would be wrong for me to force my students to do what I tell

---

<sup>27</sup> Where both of these statements are true, we have an instance of what Jeremy Waldron calls a “right to do wrong” (1981). I agree with Waldron that instances in which people have such a right are a familiar part of social and moral life.

them to. Legitimate authority, then, need not be a threat to people's right to autonomy or non-interference, if understood in this sense.

There is, alternatively, a weaker interpretation of what a right to autonomy or non-interference amounts to. As already noted, all legitimate authority poses a threat to people's autonomy in the following sense: if someone else has legitimate authority over me, I am no longer, normatively speaking, free to act on my own best judgment. We might, then, understand a "right to autonomy" or a "right to decide" in terms of a right to remain free from this kind of normative interference. A right to decide, in this sense, would involve a right to remain normatively free to make our own judgments, or a right to have our own normative situation unchanged by others' directives. If we had such a right, it could be used to define the contours of our legitimate sphere of control, within which no practical authority could ever be justified. The difficulty with this approach, however, is twofold. First, we might be sceptical of the existence of such a right. If we think this right exists because it protects some interest of ours, the interest in question would appear to be a normative interest, and more specifically, an interest in remaining normatively unaffected by others' commands. Why think we have such a normative interest, which goes above and beyond our more familiar interest in autonomy?<sup>28</sup>

---

<sup>28</sup> One possible, quasi-Kantian way of motivating this idea, however, was suggested to me by an anonymous reviewer. It goes as follows. We may have a right to make certain choices free from the wants, choices, or needs of others because we should be independent from others. And, the thought goes, we cannot be independent from others if we lack any domain where only our own judgments are sovereign. 'Independence,' in this sense, is a relational ideal of not being always normatively subject to the potential choices of others. I think this is an interesting suggestion. More, of course, would need to be said about what it means to be independent from others in this sense and why it is valuable. As already noted, however, I think it would be possible to spell out an alternative version of the 'independence condition' grounded in this ideal, rather than in Raz's ideal of autonomy. Either way of spelling out the independence condition would yield a sphere of action in which we are not normatively subject to the potential choices of others. On Raz's account, this is the sphere of action in which it is more important to make our own decisions, than to get things right. On the alternative, this would presumably be the sphere of action which is necessary to guarantee our 'independence.' Importantly, however, both accounts would remain compatible with an instrumentalist justification of authority relations. Further, for those who are attracted by this ideal of independence, it is worth noting that Raz's account of autonomy appears to already include an appeal to something like it. In the *Morality of Freedom*, Raz takes 'independence' to be one of three distinct components of the conditions of autonomy, alongside appropriate mental abilities and an adequate range of options (1986, p. 372). As

Furthermore, even if we accept the existence of such a normative interest, we would have to specify which spheres of action such a right would extend to. It would be implausible to posit a universal right to having our normative situation unchanged by others' directives. If we did, it would rule out the possibility of justified practical authority altogether. So, in which spheres of action do we have a right to have our normative situation unchanged by others' directives? One might be tempted to answer: "the sphere of action which concerns only myself, my own goals and projects." But as I've showed in this article, if we accept this answer, we would rule out the possibility of legitimate practical authority in a range of central and paradigmatic cases of practical authority, such as cases involving the authority of parents and teachers.

It seems to me, then, that any concern about paternalism which is motivated by either an outcome-based, or a motive-based account, is best accommodated by a condition like the independence condition, in combination with a suitably robust account of the value of autonomy. I think we should acknowledge, however, that we may have a very different sort of complaint against John's or Ruth's authority, in line with the sort of complaints against paternalism that "expressive-content" accounts describe. Consider the fact that, even if their directives would give me protected reasons for action, there may be very good reasons why John or Ruth shouldn't issue commands to me in the first place. Proponents of "expressive-content" views of paternalism tend to point to the fact that commands might express the idea that the commander knows better than the commanded, and this can (in certain circumstances) be disrespectful – even where it is true that the commander knows better than the commanded. To illustrate, it would be disrespectful of me to point out to a stranger on the subway that their shoes look cheap, even if this is true, and open for all to see. It seems right to say that we might have a paternalism-related reason to complain against John or Ruth issuing authoritative directives to us, precisely because of what such commands might express.<sup>29</sup>

But if this is the kind of paternalism we are worried about, it should be clear, first, that narrowing the scope of instrumentalist accounts of authority would not help to discharge the

---

Raz uses the term, independence is about making decisions while remaining uninfluenced by another's will.

<sup>29</sup> Note, though, that many instances of advice-giving have the same expressive content as command-giving: namely, that the person who gives advice knows better than the person to whom advice is given. While expressive-content views of paternalism thus capture one aspect of what we sometimes find objectionable about relations of command and obedience, they do not capture the special moral problem of command as distinguished from counsel.

worry. People have moral obligations of many different kinds: obligations to keep the promises they make to their friends, to refrain from revealing their colleagues' secrets, to take care of their family members. One might be suspicious of a state which attempts to give people commands aimed at helping them improve their compliance with pre-existing moral obligations of this kind. Such directives could be criticised on the grounds that they would be objectionably paternalistic, precisely because they express the kind of negative or disrespectful attitude about citizens' ability to run their own lives that proponents of "expressive-content" views – including Quong himself – think lies at the heart of objectionable forms of paternalism.

More importantly, however, it is not clear why, even if John's or Ruth's directives were objectionably paternalistic in this sense, it would undermine their claim to practical authority over me. The sense of authority that I have been concerned with, throughout this piece, is the ability to impose obligations on others. Why would the fact that John or Ruth has good reason not to issue a certain directive to me undermine the bindingness of that directive, once issued? Compare: the fact that it would violate the norms of friendship if I started issuing commands to a friend – which, if obeyed, would be genuinely helpful for her – certainly gives *me* a good reason not to issue commands to her. If I proceed anyways, this may give my friend a good reason to take offence or get mad; if I do so repeatedly, it may give her a good reason to end our friendship. But does it also defeat the reason my friend has to do as I told her to? I do not think it does. The same point holds also in the case paternalistic advice. I may resent you for giving me paternalistic advice, and you may have good reason to avoid giving me such advice. That doesn't imply that I shouldn't do what you advised me to, given your advice was good advice. Paternalistic commands, much like paternalistic advice, can still be reason-giving.

We should acknowledge, however, that these considerations might lead us to want to distinguish between authority as I have understood the term here, in terms of the *ability* to issue binding directives, and authority understood in terms of the *permission*, or the *right*, to issue binding directives. (In a similar vein, others have suggested that we might helpfully distinguish between *justified* and *legitimate* authority, where justified authority refers to the ability or power to impose obligations, and legitimate authority refers to, roughly, justified authority *plus* morally acceptable relations between ruler and ruled.)<sup>30</sup> While reasons which speak against issuing authoritative directives may not undermine the bindingness of such directives, it should be clear that they could undermine the permission or right to issue them in the first place. Clearly, reasons related to paternalism can speak against issuing directives, but there are other

---

<sup>30</sup> See e.g. Viehoff (2016).

situations, too, in which the relationship between ruler and ruled is flawed in ways that may undermine the ruler's permission or right to issue binding directives. For example, if a ruler creates the very problem that their authoritative directives are meant to solve, their authoritative directives may well be binding, in spite of the fact that there is something unacceptable about the relation between ruler and ruled. As Viehoff (2016, p. 120) points out, it isn't surprising that the questions of whether one person has the ability to impose obligations on another, and the question of whether the relation between ruler and ruled is morally acceptable along various dimensions, can come apart. After all, the fact that A's power to bind B is morally unacceptable needn't undo the instrumental value of such power.

What I do want to insist on, however, is the usefulness of a relatively undemanding account of what practical authority amounts to. Justified practical authority, in the sense that I have been using the term here, involves the normative capacity or ability to impose obligations. Practical authority, however, is often thought to involve much more than this. Authority is often analysed in terms of a "right to rule," where this is understood to encompass a Hohfeldian power to change subjects' normative situation and a claim-right, on the part of authorities, to their subjects' obedience. We often associate political authority, specifically, with an even more extensive list of rights: for example, a (claim-)right to non-interference, a (power-)right to enforce, a (claim- or liberty-)right to err – that is, a permission to make certain kinds of mistakes in issuing commands, as well as the avoidance of liability for at least some of the costs that may arise as a result of such mistakes.<sup>31</sup> The idea that practical authorities enjoy a claim-right to subjects' obedience is also closely related to the thought that authorities are wronged when subjects fail to comply with their commands. But that someone has practical authority in our sense need not imply that they have any of these rights, nor that they would be wronged if the subject were to disobey their commands.<sup>32</sup>

---

<sup>31</sup> By a right to err, I mean not only an ability or permission to issue binding but mistaken directives, but also a certain form of insulation against compensatory demands that might arise from such mistakes. For a helpful discussion, see Viehoff (2019).

<sup>32</sup> This point, I think, also comes out of a careful reading of Raz. Though Raz sometimes refers to authority as a "right to rule" (e.g. 1986, p. 23) – an idea traditionally understood in terms of a *power* to impose obligations (and change subjects' normative situation in other ways) combined with a *claim-right* to subjects' obedience – his justification of authority focuses squarely on the normative power to issue binding directives, rather than on justifying any further claim-rights that some authorities may also enjoy. Furthermore, Raz is uncomfortable with the suggestion that we owe it to legitimate

Several commentators have complained that understanding practical authority on these relatively undemanding terms is revisionist, because it departs too much from our ordinary understanding of authority.<sup>33</sup> I do not think this is right. First, it is not obvious that our ordinary notions of political or parental authority involve the idea that authorities have a claim-right to subjects' obedience. Claim-rights tend to correspond to directed duties, but it isn't obvious that citizens or children owe it to their political representatives, or to their parents, to obey. Politicians and parents ought to rule on their subjects' behalf. They exercise power for the sake of their subjects' benefit. It seems more accurate, then, to suggest that we owe it to ourselves (or to our fellow citizens) to obey, if and when we do. (Other times, we might owe it to no one at all. This might be true, for example, of authoritative political directives which tell us to respect places of natural beauty.) Furthermore, it is indicative that the normal consequence of breaking the law, or disobeying one's parents, isn't restitution or repair, but punishment. And while it might, in some cases, be appropriate to apologise to one's parents for disobeying them, or to try to make it up to them if one does, this can be explained by the fact that, while authority tends to play some role in the relationship between parents and children, it by no means exhausts parent-child relationships. Such relationships involve all kinds of emotional and material dependencies, attachments, and agreements which can make apologies and amends fitting.<sup>34</sup>

Second, whether we think an account of authority is revisionist or not appears to depend on what kind of practical authority we have in mind. While it is natural to think of e.g. political authority as involving an extensive list of rights, going above and beyond the normative power to issue obligations, it is not obvious that the practical authority we associate with teachers, bosses or sports coaches involves any such rights. If we are trying to develop and defend a unified account of practical authority, then, it doesn't look revisionist to focus solely on the normative capacity or power to impose obligations understood in terms of protected reasons. On the contrary, a more demanding account of what having practical authority amounts to

---

authorities to obey them. And in response to Darwall's comment that he cannot explain subjects' accountability to authorities for doing as commanded, Raz replies that he is “more used to the idea that those in authority are accountable to their subjects than to the thought that their subjects are accountable to them” (Raz 2010, p. 299).

<sup>33</sup> See e.g. Hershovitz (2011, p. 10-11) for this complaint.

<sup>34</sup> For two more detailed argument against the idea that authority should be understood in terms of a right to rule, see Marmor (2011, p. 255-260) and Enoch (2014, p. 323-328).

would risk being more radically revisionist, since it threatens to exclude many familiar cases of practical authority altogether.

Let's return, then, to Darwall's counterexample, keeping in mind what legitimate practical authority, as I am understanding it here, amounts to. First and most obviously, to say that Ruth has legitimate authority in this sense does not imply anything about the permissibility of enforcement, or the use of coercion. Directives can be legitimate, even when it would be wrong to enforce them. Second, when claiming that someone like Ruth might have legitimate practical authority over me, we do not have to think that she has any right to my obedience, nor that she would be wronged if I disobeyed her. Finally, in so far as I can drop or change my goals or projects, there is a sense in which the directives of Ruth remain optional for me. I insist, then, that we should not be worried about granting that Ruth has legitimate practical authority over me.

## 5. Conclusion

I've argued that instrumentalist accounts of authority would not be more plausible if restricted in scope to some particularly weighty subset of moral or categorical concerns; and that the widespread assumption that legitimate authoritative directives must have the force of moral or categorical reasons is mistaken. Furthermore, I've suggested that worries about paternalism do not pose any particular threat to instrumentalist accounts of authority.

The approach I've taken naturally gives rise to the following worry. Why think it possible to give a unified account of practical authority in the first place? Or alternatively, even if this is possible, why think such an approach to analysing practical authority is more useful or valuable than an approach which looks at the different kinds of practical authority separately, asking not what justifies practical authority in general, but what justifies specific kinds of practical authority, in specific contexts? The best response to this worry, I think, is to point out that there is clearly something that we take to be recognisably similar about the practical authority of politicians and judges, parents and officers, teachers and coaches, and it seems like a worthwhile philosophical pursuit to try to understand what this something consists of. Furthermore, the fact that we may have to complement this general story about practical authority with a story about what it takes to justify the more demanding forms of practical authority that we find in certain contexts does not speak against the usefulness or value of such a general account. I for one, do not think that Raz's service conception of authority can explain why e.g. political authorities might have a right to coercively enforce their directives, nor why

they should enjoy a right against interference, or be able to avoid liability for certain kinds of costs related to their issuing and enforcing mistaken or unjust directives. But nothing I have said here precludes us from giving a further story which could explain these features of specifically political authority.

I take my argument to have shown, however, that those who think (for example) that political authority is illegitimate unless based on a certain particularly weighty kind of moral reason, cannot use an argument about the nature of legitimate practical authority *in general* to support their position. Rather, what they need is an argument which shows political authority to be fundamentally different from other familiar kinds of practical authority.

## References

- Christiano, Thomas (2004). “The authority of democracy.” *Journal of Political Philosophy* 12 (3):266–290.
- Coleman, Jules L. (2009). “Beyond inclusive legal positivism.” *Ratio Juris* 22 (3):359-394.
- Cornell, Nicolas (2015). A Third Theory of Paternalism. *Michigan Law Review* 113:1295-1336.
- Darwall, Stephen (2013). *Morality, Authority, and Law: Essays in Second-Personal Ethics I*. Oxford University Press.
- Enoch, David (2014). “Authority and Reason-Giving.” *Philosophy and Phenomenological Research* 89 (2):296-332.
- Estlund, David (2009). *Democratic Authority: A Philosophical Framework*. Princeton University Press.
- Hart, H. L. A. (1982). *Essays on Bentham: Jurisprudence and Political Philosophy*. Oxford, GB: Oxford University Press.
- Hershovitz, Scott (2011). “The Role of Authority.” *Philosophers' Imprint* 11(7):1-19.
- Hobbes, Thomas and Gaskin, J. C. A. (ed.) (2008). *Leviathan*. Oxford University Press.
- Marmor, Andrei (2011). “An Institutional Conception of Authority.” *Philosophy and Public Affairs* 39 (3):238-261.
- Monti, Ezequiel H. (2018). “On Darwall’s Case against the Normal Justification Thesis.” *Ethics* 128 (2):432-445.
- Perry, Stephen R. (2013). “Political Authority and Political Obligation.” In *Oxford Studies in Philosophy of Law: Volume 2*. Oxford: Oxford University Press.

Quong, Jonathan (2010). *Liberalism Without Perfection*. Oxford University Press.

Rawls, John (1993). *Political Liberalism*. Columbia University Press.

Raz, Joseph (1977). “Promises and Obligations.” In *Law, Morality, and Society: Essays in Honour of H. L. A. Hart*, ed. P. M. S. Hacker and Joseph Raz. Oxford: Oxford University Press.

Raz, Joseph (1979). *The Authority of Law: Essays on Law and Morality*. Oxford University Press.

Raz, Joseph (1986). *The Morality of Freedom*. Oxford University Press.

Raz, Joseph (1989). “Facing Up: A Reply.” *Southern California Law Review* 62:1153–236.

Raz, Joseph (2001). *Value, Respect, and Attachment*. Cambridge University Press.

Raz, Joseph (2009). *Between Authority and Interpretation*. Oxford University Press.

Raz, Joseph (2010). ”On respect, authority, and neutrality: A response.” *Ethics* 120 (2):279-301.

Raz, Joseph (2011). *From Normativity to Responsibility*. Oxford University Press.

Raz, Joseph (2022). *The Roots of Normativity*. Ed. Ulrike Heuer. Oxford: Oxford University Press.

Shiffrin, Seana Valentine (2000). Paternalism, Unconscionability Doctrine, and Accommodation. *Philosophy and Public Affairs* 29 (3):205-250.

Sunstein, Cass R. and Thaler, Richard H. (2003) “Libertarian Paternalism Is Not an Oxymoron,” *University of Chicago Law Review*: Vol. 70: Iss. 4, Article 1.

Tadros, Victor (2020). *To Do, to Die, to Reason Why: Individual Ethics in War*. Oxford University Press.

Viehoff, Daniel (2011). “Debate: Procedure and Outcome in the Justification of Authority.” *Journal of Political Philosophy* 19 (2):248-259.

Viehoff, Daniel (2016). “On Authority, Legitimacy, and Service: Discussion of The Morality of Freedom.” *Jerusalem Review of Legal Studies* 14(1):113-124.

Viehoff, Daniel (2019). “Legitimacy as a Right to Err.” In Jack Knight & Melissa Schwartzberg (eds.), *NOMOS LXI: Political Legitimacy*. New York: NYU Press.

Viehoff, Daniel (forthcoming). *Serving the Governed*.

Waldron, Jeremy (1981). A right to do wrong. *Ethics* 92 (1):21-39.